

Implementing Basic Vector Search

A Practical Guide to Building Your First Qdrant Collection

In this guide, you will build:

1. A Connection to Qdrant Cloud
2. A Vector Collection
3. A Similarity Search Engine

Created by Raj Ghosh

Contents

1	Environment Setup	2
1.1	Installing and Importing	2
1.2	Connecting to the Cloud	2
2	The Collection (Your Database Table)	3
2.1	Understanding Collections	3
3	Inserting Points	4
3.1	Anatomy of a Point	4
4	Running Similarity Search	5

Chapter 1

Environment Setup

1.1 Installing and Importing

To speak to the Qdrant database, we need the official Python translator (the Client).

Step 1: Install Client

Run this command in your terminal or notebook to install the library.

```
1 !pip install qdrant-client
```

Step 2: Import Libraries

We need to bring in two key components:

- `QdrantClient`: The "phone" we use to call the database.
- `models`: The "dictionary" containing definitions for Points, Vectors, and Distances.

```
1 from qdrant_client import QdrantClient, models  
2 import os
```

1.2 Connecting to the Cloud

Step 3: Connect

Using the API key and URL you retrieved in Day 0, establish the link.

```
1 # Load secrets from environment (Best Practice)  
2 qdrant_url = os.getenv("QDRANT_URL")  
3 qdrant_key = os.getenv("QDRANT_API_KEY")  
4  
5 client = QdrantClient(url=qdrant_url, api_key=qdrant_key)  
6  
7 # PRO TIP: For quick testing without internet, use Memory mode:  
8 # client = QdrantClient(":memory:")
```

Chapter 2

The Collection (Your Database Table)

2.1 Understanding Collections


A **Collection** in Qdrant is equivalent to a Table in SQL. It holds your vectors. When you make one, you must define the "Rules of the Game" (Vector Size and Distance Metric).

Step 4: Create Collection

```
1 collection_name = "my_first_collection"
2
3 client.create_collection(
4     collection_name=collection_name,
5     vectors_config=models.VectorParams(
6         size=4,                                     # We are using 4D vectors
7         distance=models.Distance.COSINE           # How to measure similarity
8     )
9 )
```

💡 Deep Dive: Distance Metrics Explained

How do we know if Vector A is "close" to Vector B?



Step 5: Verify Creation

Always double-check that the operation succeeded.

```
1 collections = client.get_collections()
2 print(collections)
3 # Output should contain "my_first_collection"
```

Chapter 3

Inserting Points

3.1 Anatomy of a Point

In Qdrant, a data entry is called a **Point**.

- **ID:** Unique Number or UUID.
- **Vector:** The list of numbers (coordinates).
- **Payload:** JSON data (metadata like names, dates, categories).

Step 6: Upsert (Update/Insert)

We use ‘upsert’ to put data in. If the ID exists, it updates it. If not, it creates it.

```
1 points = [
2     models.PointStruct(
3         id=1,
4         vector=[0.1, 0.2, 0.3, 0.4],
5         payload={"category": "example"}
6     ),
7     models.PointStruct(
8         id=2,
9         vector=[0.2, 0.3, 0.4, 0.5],
10        payload={"category": "demo"}
11    )
12]
13
14 client.upsert(
15     collection_name=collection_name,
16     points=points
17)
```

Step 7: Collection Info


Verify the data is actually there.

```
1 info = client.get_collection(collection_name)
2 print(f"Points Count: {info.points_count}")
3 # Should print: Points Count: 2
```

Chapter 4

Running Similarity Search

This is the moment of truth. We provide a **Query Vector**, and Qdrant finds the mathematically closest match.



Semantic Search

Finds the vector with the highest similarity score.

Step 8: Query Points

```
1 query_vector = [0.08, 0.14, 0.33, 0.28]
2
3 search_results = client.query_points(
4     collection_name=collection_name,
5     query=query_vector,
6     limit=1  # How many results do you want?
7 )
8
9 print(search_results)
```

Understanding the Output

You will receive a `ScoredPoint` object:

- **ID:** The ID of the matching point (e.g., 1).
- **Score:** A number between 0 and 1 (for Cosine). Closer to 1.0 means more similar.
- **Payload:** The metadata we stored ("category": "example").

Tutorial Completed

You have successfully implemented your first Vector Search engine.

Created by Raj Ghosh