

# Qdrant Vector Database

Embeddings, Payloads, and Advanced Filtering

---

## 1 Common Embedding Sources

Choosing the right embedding source is a critical decision that balances **cost, performance, and accuracy**.

### 1.1 1. On-Premise, Optimized: FastEmbed by Qdrant

FastEmbed is optimized for on-premise, high-speed generation with minimal dependencies. It uses quantized weights and ONNX Runtime, making it up to **50% faster** than traditional PyTorch models.

When to choose FastEmbed:

- On-premise execution for privacy-sensitive applications.
- High-speed CPU inference without heavy dependencies.
- Tightly integrated with Qdrant.

#### FastEmbed Quickstart

```
1 from qdrant_client import QdrantClient
2 from fastembed import TextEmbedding
3
4 # Default model: BAAI/bge-small-en-v1.5 (~67MB)
5 embedding_model = TextEmbedding()
6 vector = embedding_model.embed("Qdrant is a vector search engine")
```

### 1.2 2. Managed and Integrated: Cloud Providers

- **Qdrant Cloud Inference:** Managed service directly within your cluster. Eliminates external network latency.
- **Third-Party APIs:** Commercial APIs (OpenAI, Anthropic) offering state-of-the-art models (Trade-off: network latency and costs).

### 1.3 3. On-Premise, Customizable: Open Source Models

Libraries like **Sentence Transformers** give access to the Hugging Face Hub. Ideal for fine-tuning on domain-specific data or running on local GPUs.

---

2codebgwhite

primary

---

<b>Execution</b>	On-premise (CPU/GPU)	Cloud API	On-premise (CPU/GPU)
<b>Speed Control</b>	Optimized CPU latency High	API latency Low	Varies by hardware Maximum

---

## 2 Payloads (Metadata)

Vectors capture semantic relevance, but **Payloads hold structured metadata** for business logic filtering.

### 2.1 Supported Payload Types

- **Keyword:** Exact string matching (`category: "electronics"`).
- **Integer / Float:** Numerical filtering (`price: 19.99`).
- **Bool:** True/false values.
- **Geo:** Latitude/longitude maps.
- **Datetime:** RFC 3339 format timestamps.
- **UUID:** Memory-efficient ID matching.


### 2.2 Data Structures

Payloads can store complex structures:

- **Arrays:** Succeeds if *at least one* value matches (`tags: ["vegan", "organic"]`).
- **Nested Objects:** Queried using dot notation (`user.address.city`).

## 3 Filtering Logic: Building Complex Queries

### 3.1 Logical Clauses



Complex Filter Example

```
1 models.Filter(  
2     should=[  
3         models.Filter(must=[  
4             models.FieldCondition(key="category", match=models.MatchValue(  
5                 value="electronics")),  
6             models.FieldCondition(key="price", range=models.Range(lt=200))  
7         ]),  
8         models.Filter(must=[  
9             models.FieldCondition(key="category", match=models.MatchValue(  
10                value="books")),  
11             models.FieldCondition(key="rating", range=models.Range(gte=4.0))  
12         ])  
13     ]  
14 )
```

## 3.2 Advanced Filtering: Nested Objects

Nested filtering ensures conditions are evaluated within *individual* array elements rather than across all elements.

Nested Condition Example

```
1 models.Filter(
2     must=[
3         models.NestedCondition(
4             nested=models.Nested(
5                 key="reviews",
6                 filter=models.Filter(must=[
7                     models.FieldCondition(key="rating", match=models.
8                         MatchValue(value=5)),
9                     models.FieldCondition(key="verified", match=models.
10                        MatchValue(value=True))
11                 ]))
12             )
13         ])
```

## 4 Performance Optimization

Create payload indexes for frequently filtered fields to maximize performance.

### Optimization Note

When filters are highly selective, Qdrant's query planner may bypass vector indexing entirely and use payload indexes for faster results.

Creating Indexes

```
1 # Index frequently filtered fields
2 client.create_payload_index(
3     collection_name="{collection_name}",
4     field_name="category",
5     field_schema=models.PayloadSchemaType.KEYWORD,
6 )
7
8 # For multi-tenant applications
9 client.create_payload_index(
10    collection_name="{collection_name}",
11    field_name="tenant_id",
12    field_schema=models.KeywordIndexParams(type="keyword", is_tenant=True),
13 )
```

## 5 Key Takeaways

- 1. Flexibility:** Points combine unique IDs, vectors, and structured metadata.
- 2. Embedding Strategy:** Balance speed, accuracy, and privacy using FastEmbed, Cloud APIs, or Open Source models.
- 3. Powerful Filtering:** Use `must`, `should`, and `must_not` alongside specialized payload indexes to bridge semantic search with strict business rules.